

dwjesr11i

September 11, 2023

```
[5]: # Linear algebra
import numpy as np

# Data manipulation and analysis
import pandas as pd

# Data visualization
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from matplotlib import style

# Algorithms
from sklearn import linear_model
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
```

```
[7]: train_df = pd.read_csv('taks_2 Train.csv')
test_df = pd.read_csv('taks_2 Test.csv')

train_df['train_test'] = 1
test_df['train_test'] = 0
# test_df['Survived'] = np.NaN
all_data = pd.concat([train_df, test_df])

%matplotlib inline
all_data.columns
```

```
[7]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
        'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked', 'train_test'],
        dtype='object')
```

```
[8]: train_df.head(10)
```

```
[8]:   PassengerId  Survived  Pclass \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
5             6         0        3
6             7         0        1
7             8         0        3
8             9         1        3
9            10         1        2
```

```

                                Name      Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                        Heikkinen, Miss. Laina  female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4                        Allen, Mr. William Henry    male  35.0      0
5                        Moran, Mr. James    male   NaN      0
6                        McCarthy, Mr. Timothy J    male  54.0      0
7                        Palsson, Master. Gosta Leonard    male   2.0      3
8  Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)  female  27.0      0
9                        Nasser, Mrs. Nicholas (Adele Achem)  female  14.0      1

```

```

   Parch      Ticket    Fare Cabin Embarked  train_test
0      0    A/5 21171   7.2500   NaN        S          1
1      0    PC 17599  71.2833   C85        C          1
2      0  STON/O2. 3101282   7.9250   NaN        S          1
3      0    113803  53.1000  C123        S          1
4      0    373450   8.0500   NaN        S          1
5      0    330877   8.4583   NaN        Q          1
6      0    17463  51.8625   E46        S          1
7      1    349909  21.0750   NaN        S          1
8      2    347742  11.1333   NaN        S          1
9      0    237736  30.0708   NaN        C          1

```

```
[9]: test_df.head(10)
```

```
[9]:   PassengerId  Pclass                                Name      Sex  \
0           892        3                        Kelly, Mr. James    male
1           893        3      Wilkes, Mrs. James (Ellen Needs)  female
2           894        2            Myles, Mr. Thomas Francis    male
3           895        3                        Wirz, Mr. Albert    male
4           896        3  Hirvonen, Mrs. Alexander (Helga E Lindqvist)  female
5           897        3      Svensson, Mr. Johan Cervin    male
```

6	898	3	Connolly, Miss. Kate	female
7	899	2	Caldwell, Mr. Albert Francis	male
8	900	3	Abraham, Mrs. Joseph (Sophie Halaut Easu)	female
9	901	3	Davies, Mr. John Samuel	male

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	train_test
0	34.5	0	0	330911	7.8292	NaN	Q	0
1	47.0	1	0	363272	7.0000	NaN	S	0
2	62.0	0	0	240276	9.6875	NaN	Q	0
3	27.0	0	0	315154	8.6625	NaN	S	0
4	22.0	1	1	3101298	12.2875	NaN	S	0
5	14.0	0	0	7538	9.2250	NaN	S	0
6	30.0	0	0	330972	7.6292	NaN	Q	0
7	26.0	1	1	248738	29.0000	NaN	S	0
8	18.0	0	0	2657	7.2292	NaN	C	0
9	21.0	2	0	A/4 48871	24.1500	NaN	S	0

```
[10]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age            714 non-null    float64
6   SibSp           891 non-null    int64
7   Parch           891 non-null    int64
8   Ticket          891 non-null    object
9   Fare            891 non-null    float64
10  Cabin           204 non-null    object
11  Embarked        889 non-null    object
12  train_test      891 non-null    int64
dtypes: float64(2), int64(6), object(5)
memory usage: 90.6+ KB
```

```
[11]: train_df.describe()
```

```
[11]:
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	

25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare	train_test
count	891.000000	891.000000	891.0
mean	0.381594	32.204208	1.0
std	0.806057	49.693429	0.0
min	0.000000	0.000000	1.0
25%	0.000000	7.910400	1.0
50%	0.000000	14.454200	1.0
75%	0.000000	31.000000	1.0
max	6.000000	512.329200	1.0

```
[12]: total = train_df.isnull().sum().sort_values(ascending=False)
percent_1 = train_df.isnull().sum()/train_df.isnull().count()*100
percent_2 = (round(percent_1, 1)).sort_values(ascending=False)
missing_data = pd.concat([total, percent_2], axis=1, keys=['Total', '%'])
missing_data.head(13)
```


```
[12]:
```

	Total	%
Cabin	687	77.1
Age	177	19.9
Embarked	2	0.2
PassengerId	0	0.0
Survived	0	0.0
Pclass	0	0.0
Name	0	0.0
Sex	0	0.0
SibSp	0	0.0
Parch	0	0.0
Ticket	0	0.0
Fare	0	0.0
train_test	0	0.0

```
[13]: train_df.columns.values
```

```
[13]: array(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
        'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked', 'train_test'],
        dtype=object)
```

```
[14]: sns.boxplot(x='Survived',y='Fare',data=train_df);
```



```
[15]: train_df[train_df['Fare']>300]
```

```
[15]:
```

	PassengerId	Survived	Pclass	Name \
258	259	1	1	Ward, Miss. Anna
679	680	1	1	Cardeza, Mr. Thomas Drake Martinez
737	738	1	1	Lesurer, Mr. Gustave J

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	\
258	female	35.0	0	0	PC 17755	512.3292	NaN	C	
679	male	36.0	0	1	PC 17755	512.3292	B51 B53 B55	C	
737	male	35.0	0	0	PC 17755	512.3292	B101	C	


```

train_test
258      1
679      1
737      1

```


```
[16]: train_df[train_df['Name'].str.contains("Capt")]
```

```
[16]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	\
745	746	0	1	Crosby, Capt. Edward Gifford	male	70.0	

	SibSp	Parch	Ticket	Fare	Cabin	Embarked	train_test
745	1	1	WE/P 5735	71.0	B22	S	1

```
[17]: FacetGrid = sns.FacetGrid(train_df, col='Embarked', height=4, aspect=1.2)
FacetGrid.map(sns.pointplot, 'Pclass', 'Survived', 'Sex', ci=95.0,
↪palette='deep', order=None, hue_order=None)
FacetGrid.add_legend();
```




```
[18]: sns.set(style='darkgrid')
plt.subplots(figsize = (8,6))
ax=sns.countplot(x='Sex', data = train_df, hue='Survived', edgecolor=(0,0,0),
↪linewidth=2)

# Fixing title, xlabel and ylabel
plt.title('Passenger distribution of survived vs not-survived', fontsize=25)
plt.xlabel('Gender', fontsize=15)
plt.ylabel("# of Passenger Survived", fontsize = 15)
labels = ['Female', 'Male']

# Fixing xticks.
plt.xticks(sorted(train_df.Survived.unique()),labels);
```

Passenger distribution of survived vs not-survived



```
[19]: train_df.groupby(['Sex']).mean()
```

```
[19]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	\
Sex							
female	431.028662	0.742038	2.159236	27.915709	0.694268	0.649682	
male	454.147314	0.188908	2.389948	30.726645	0.429809	0.235702	


```
Fare train_test
```

Sex	Fare	train_test
female	44.479818	1.0
male	25.523893	1.0

```
[20]: train_df.groupby(['Sex', 'Pclass']).mean()
```

```
[20]:
```

		PassengerId	Survived	Age	SibSp	Parch	\
Sex	Pclass						
female	1	469.212766	0.968085	34.611765	0.553191	0.457447	
	2	443.105263	0.921053	28.722973	0.486842	0.605263	
	3	399.729167	0.500000	21.750000	0.895833	0.798611	
male	1	455.729508	0.368852	41.281386	0.311475	0.278689	
	2	447.962963	0.157407	30.740707	0.342593	0.222222	

3 455.515850 0.135447 26.507589 0.498559 0.224784

	Sex	Pclass	Fare	train_test
female	1	106.125798	1.0	
	2	21.970121	1.0	
	3	16.118810	1.0	
male	1	67.226127	1.0	
	2	19.741782	1.0	
	3	12.661633	1.0	

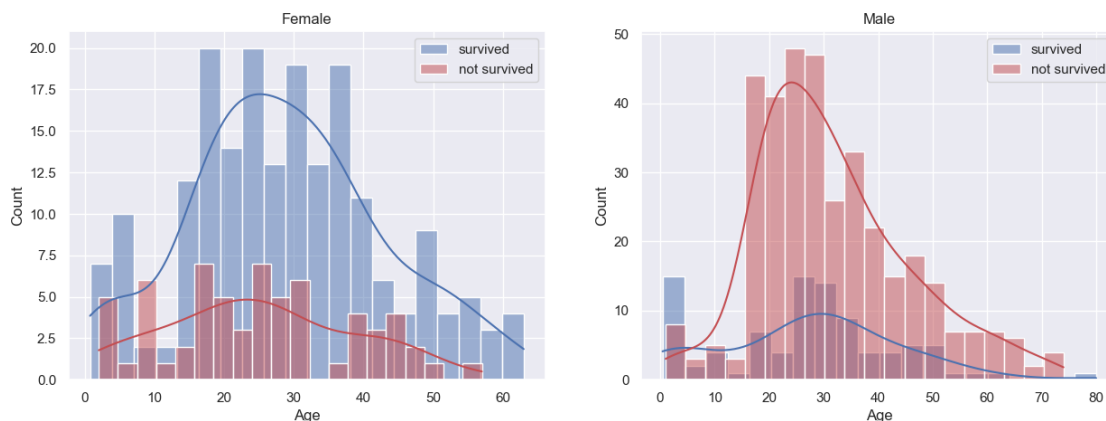
```
[21]: survived = 'survived'
not_survived = 'not survived'

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(15, 5))

women = train_df[train_df['Sex']=='female']
men = train_df[train_df['Sex']=='male']

# Plot Female Survived vs Not-Survived distribution
ax = sns.histplot(women[women['Survived']==1].Age.dropna(), bins=20, label = 'survived', ax = axes[0], color='b', kde=True)
ax = sns.histplot(women[women['Survived']==0].Age.dropna(), bins=20, label = 'not_survived', ax = axes[0], color='r', kde=True)
ax.legend()
ax.set_title('Female')

# Plot Male Survived vs Not-Survived distribution
ax = sns.histplot(men[men['Survived']==1].Age.dropna(), bins=20, label = 'survived', ax = axes[1], color='b', kde=True)
ax = sns.histplot(men[men['Survived']==0].Age.dropna(), bins=20, label = 'not_survived', ax = axes[1], color='r', kde=True)
ax.legend()
ax.set_title('Male');
```




```
[22]: train_df[train_df['Age']<18].groupby(['Sex','Pclass']).mean()
```


```
[22]:
```

		PassengerId	Survived	Age	SibSp	Parch	\
female	1	525.375000	0.875000	14.125000	0.500000	0.875000	
	2	369.250000	1.000000	8.333333	0.583333	1.083333	
	3	374.942857	0.542857	8.428571	1.571429	1.057143	
male	1	526.500000	1.000000	8.230000	0.500000	2.000000	
	2	527.818182	0.818182	4.757273	0.727273	1.000000	
	3	437.953488	0.232558	9.963256	2.069767	1.000000	

		Fare	train_test
female	1	104.083337	1.0
	2	26.241667	1.0
	3	18.727977	1.0
male	1	116.072900	1.0
	2	25.659473	1.0
	3	22.752523	1.0

```
[23]: plt.subplots(figsize = (8,8))
ax=sns.countplot(x='Pclass',hue='Survived',data=train_df)
plt.title("Passenger Class Distributon - Survived vs Non-Survived", fontsize = 25);
```


Passenger Class Distribution - Survived vs Non-Survived



```
[24]: plt.subplots(figsize=(10,8))
ax=sns.kdeplot(train_df.loc[(train_df['Survived'] == 0), 'Pclass'],shade=True,color='r',label='Not Survived')
ax.legend()
ax=sns.kdeplot(train_df.loc[(train_df['Survived'] == 1), 'Pclass'],shade=True,color='b',label='Survived')
ax.legend()


plt.title("Passenger Class Distribution - Survived vs Non-Survived", fontsize = 25)
labels = ['First', 'Second', 'Third']
plt.xticks(sorted(train_df.Pclass.unique()),labels);
```

Passenger Class Distribution - Survived vs Non-Survived




```
[25]: plt.subplots(figsize = (8,6))
sns.barplot(x='Pclass', y='Survived', data=train_df);
plt.title("Passenger Class Distribution - Survived Passengers", fontsize = 25);
```

Passenger Class Distribution - Survived Passengers



```
[26]: # Look at numeric and categorical values separately
df_num = train_df[['Age', 'SibSp', 'Parch', 'Fare']]
df_cat = train_df[['Survived', 'Pclass', 'Sex', 'Ticket', 'Cabin', 'Embarked']]
```

```
[27]: plt.subplots(figsize = (12,6))
sns.heatmap(df_num.corr(), annot=True, cmap="RdBu")
plt.title("Correlations Among Numeric Features", fontsize = 18);
```



```
[28]: train_df = train_df.drop(['PassengerId'], axis=1)
      train_df.head()
```

```
[28]:
```

	Survived	Pclass	Name \
0	0	3	Braund, Mr. Owen Harris
1	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...
2	1	3	Heikkinen, Miss. Laina
3	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)
4	0	3	Allen, Mr. William Henry


	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked \
0	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	female	38.0	1	0	PC 17599	71.2833	C85	C
2	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	female	35.0	1	0	113803	53.1000	C123	S
4	male	35.0	0	0	373450	8.0500	NaN	S

```
train_test
0      1
1      1
2      1
3      1
4      1
```

```
[29]: data = [train_df, test_df]
for dataset in data:
    dataset['relatives'] = dataset['SibSp'] + dataset['Parch']
    dataset.loc[dataset['relatives'] > 0, 'not_alone'] = 0
    dataset.loc[dataset['relatives'] == 0, 'not_alone'] = 1
    dataset['not_alone'] = dataset['not_alone'].astype(int)
train_df['not_alone'].value_counts()
```

```
[29]: 1    537
      0    354
      Name: not_alone, dtype: int64
```

```
[30]: plt.subplots(figsize = (16,4))
ax = sns.lineplot(x='relatives',y='Survived', data=train_df)
```



```
[31]: import re
deck = {"A": 1, "B": 2, "C": 3, "D": 4, "E": 5, "F": 6, "G": 7, "U": 8}
data = [train_df, test_df]

for dataset in data:
    dataset['Cabin'] = dataset['Cabin'].fillna("U0")
    dataset['Deck'] = dataset['Cabin'].map(lambda x: re.compile("([a-zA-Z]+)").
    ↪search(x).group())
    dataset['Deck'] = dataset['Deck'].map(deck)
    dataset['Deck'] = dataset['Deck'].fillna(0)
    dataset['Deck'] = dataset['Deck'].astype(int)
```

```
[32]: # We can now drop the Cabin feature
train_df = train_df.drop(['Cabin'], axis=1)
test_df = test_df.drop(['Cabin'], axis=1)
```

```
[33]: data = [train_df, test_df]

for dataset in data:
    mean = train_df["Age"].mean()
```

```

std = test_df["Age"].std()
is_null = dataset["Age"].isnull().sum()

# Compute random numbers between the mean, std and is_null
rand_age = np.random.randint(mean - std, mean + std, size = is_null)

# Fill NaN values in Age column with random values generated
age_slice = dataset["Age"].copy()
age_slice[np.isnan(age_slice)] = rand_age
dataset["Age"] = age_slice
dataset["Age"] = train_df["Age"].astype(int)

```

```
[34]: train_df["Age"].isnull().sum()
```

```
[34]: 0
```

```
[35]: train_df['Embarked'].describe()
```

```

[35]: count      889
      unique       3
      top         S
      freq       644
      Name: Embarked, dtype: object

```

```

[36]: common_value = 'S'
      data = [train_df, test_df]

      for dataset in data:
          dataset['Embarked'] = dataset['Embarked'].fillna(common_value)

```

```
[37]: train_df['Embarked'].isnull().sum()
```

```
[37]: 0
```

```
[38]: train_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Survived    891 non-null    int64
 1   Pclass      891 non-null    int64
 2   Name        891 non-null    object
 3   Sex         891 non-null    object
 4   Age         891 non-null    int32
 5   SibSp       891 non-null    int64

```

```

6   Parch      891 non-null   int64
7   Ticket     891 non-null   object
8   Fare       891 non-null   float64
9   Embarked   891 non-null   object
10  train_test  891 non-null   int64
11  relatives  891 non-null   int64
12  not_alone  891 non-null   int32
13  Deck       891 non-null   int32
dtypes: float64(1), int32(3), int64(6), object(4)
memory usage: 87.1+ KB

```

```

[39]: data = [train_df, test_df]

for dataset in data:
    dataset['Fare'] = dataset['Fare'].fillna(0)
    dataset['Fare'] = dataset['Fare'].astype(int)

```

```

[40]: train_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null   int64
1   Pclass      891 non-null   int64
2   Name        891 non-null   object
3   Sex         891 non-null   object
4   Age         891 non-null   int32
5   SibSp       891 non-null   int64
6   Parch       891 non-null   int64
7   Ticket      891 non-null   object
8   Fare        891 non-null   int32
9   Embarked    891 non-null   object
10  train_test  891 non-null   int64
11  relatives   891 non-null   int64
12  not_alone   891 non-null   int32
13  Deck        891 non-null   int32
dtypes: int32(4), int64(6), object(4)
memory usage: 83.7+ KB

```

```

[41]: data = [train_df, test_df]
titles = {"Mr": 1, "Miss": 2, "Mrs": 3, "Master": 4, "Other": 5}

for dataset in data:
    # Extract titles
    dataset['Title'] = dataset.Name.str.extract('([A-Za-z]+)\.', expand=False)

```



```

# Replace titles with a more common title or as Other
dataset['Title'] = dataset['Title'].replace(['Lady', 'Countess', 'Capt', 'Col', 'Don', 'Dr', 'Major', 'Rev', 'Sir', 'Jonkheer', 'Dona'], 'Other')
dataset['Title'] = dataset['Title'].replace('Mlle', 'Miss')
dataset['Title'] = dataset['Title'].replace('Ms', 'Miss')
dataset['Title'] = dataset['Title'].replace('Mme', 'Mrs')

# Convert titles into numbers
dataset['Title'] = dataset['Title'].map(titles)

# Filling NaN with 0 just to be safe
dataset['Title'] = dataset['Title'].fillna(0)

```

```

[42]: train_df = train_df.drop(['Name'], axis=1)
test_df = test_df.drop(['Name'], axis=1)

```

```

[43]: # Checking results
train_df.head()

```

```

[43]:
  Survived  Pclass    Sex  Age  SibSp  Parch    Ticket   Fare  \
0         0      3  male   22     1     0   A/5 21171     7
1         1      1 female   38     1     0   PC 17599    71
2         1      3 female   26     0     0  STON/O2. 3101282    7
3         1      1 female   35     1     0    113803    53
4         0      3  male   35     0     0   373450     8

  Embarked  train_test  relatives  not_alone  Deck  Title
0         S          1          1          0     8     1
1         C          1          1          0     3     3
2         S          1          0          1     8     2
3         S          1          1          0     3     3
4         S          1          0          1     8     1

```

```

[44]: genders = {"male": 0, "female": 1}
data = [train_df, test_df]

for dataset in data:
    dataset['Sex'] = dataset['Sex'].map(genders)

```

```

[45]: train_df['Ticket'].describe()

```

```

[45]: count      891
unique      681
top      347082
freq         7
Name: Ticket, dtype: object

```

```
[46]: train_df = train_df.drop(['Ticket'], axis=1)
      test_df = test_df.drop(['Ticket'], axis=1)
```

```
[47]: ports = {"S": 0, "C": 1, "Q": 2}
      data = [train_df, test_df]

      for dataset in data:
          dataset['Embarked'] = dataset['Embarked'].map(ports)
```

```
[48]: train_df.head()
```

```
[48]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	train_test \
0	0	3	0	22	1	0	7	0	1
1	1	1	1	38	1	0	71	1	1
2	1	3	1	26	0	0	7	0	1
3	1	1	1	35	1	0	53	0	1
4	0	3	0	35	0	0	8	0	1

	relatives	not_alone	Deck	Title
0	1	0	8	1
1	1	0	3	3
2	0	1	8	2
3	1	0	3	3
4	0	1	8	1

```
[49]: data = [train_df, test_df]
      for dataset in data:
          dataset['Age'] = dataset['Age'].astype(int)
          dataset.loc[ dataset['Age'] <= 11, 'Age'] = 0
          dataset.loc[(dataset['Age'] > 11) & (dataset['Age'] <= 18), 'Age'] = 1
          dataset.loc[(dataset['Age'] > 18) & (dataset['Age'] <= 22), 'Age'] = 2
          dataset.loc[(dataset['Age'] > 22) & (dataset['Age'] <= 27), 'Age'] = 3
          dataset.loc[(dataset['Age'] > 27) & (dataset['Age'] <= 33), 'Age'] = 4
          dataset.loc[(dataset['Age'] > 33) & (dataset['Age'] <= 40), 'Age'] = 5
          dataset.loc[(dataset['Age'] > 40) & (dataset['Age'] <= 66), 'Age'] = 6
          dataset.loc[ dataset['Age'] > 66, 'Age'] = 6
```

```
[50]: # Checking the distribution
      train_df['Age'].value_counts()
```

```
[50]: 4    162
      6    158
      5    147
      3    146
      2    119
      1     91
      0     68
```

Name: Age, dtype: int64

```
[51]: train_df.head()
```

```
[51]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	train_test \
0	0	3	0	2	1	0	7	0	1
1	1	1	1	5	1	0	71	1	1
2	1	3	1	3	0	0	7	0	1
3	1	1	1	5	1	0	53	0	1
4	0	3	0	5	0	0	8	0	1

	relatives	not_alone	Deck	Title
0	1	0	8	1
1	1	0	3	3
2	0	1	8	2
3	1	0	3	3
4	0	1	8	1

```
[52]: pd.qcut(train_df['Fare'], q=6)
```

```
[52]:
```

0	(-0.001, 7.0]
1	(52.0, 512.0]
2	(-0.001, 7.0]
3	(52.0, 512.0]
4	(7.0, 8.0]
...	
886	(8.0, 14.0]
887	(26.0, 52.0]
888	(14.0, 26.0]
889	(26.0, 52.0]
890	(-0.001, 7.0]

Name: Fare, Length: 891, dtype: category
Categories (6, interval[float64, right]): [(-0.001, 7.0] < (7.0, 8.0] < (8.0, 14.0] < (14.0, 26.0] < (26.0, 52.0] < (52.0, 512.0]]

```
[53]: data = [train_df, test_df]
```

```
for dataset in data:
    dataset.loc[ dataset['Fare'] <= 7, 'Fare'] = 0
    dataset.loc[(dataset['Fare'] > 7) & (dataset['Fare'] <= 8), 'Fare'] = 1
    dataset.loc[(dataset['Fare'] > 8) & (dataset['Fare'] <= 14), 'Fare'] = 2
    dataset.loc[(dataset['Fare'] > 14) & (dataset['Fare'] <= 26), 'Fare'] = 3
    dataset.loc[(dataset['Fare'] > 26) & (dataset['Fare'] <= 52), 'Fare'] = 4
    dataset.loc[dataset['Fare'] > 52, 'Fare'] = 5
    dataset['Fare'] = dataset['Fare'].astype(int)
```

```
[54]: # Checking the dataset
train_df.head(10)
```

```
[54]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	train_test	\
0	0	3	0	2	1	0	0	0	1	
1	1	1	1	5	1	0	5	1	1	
2	1	3	1	3	0	0	0	0	1	
3	1	1	1	5	1	0	5	0	1	
4	0	3	0	5	0	0	1	0	1	
5	0	3	0	4	0	0	1	2	1	
6	0	1	0	6	0	0	4	0	1	
7	0	3	0	0	3	1	3	0	1	
8	1	3	1	3	0	2	2	0	1	
9	1	2	1	1	1	0	4	1	1	

	relatives	not_alone	Deck	Title
0	1	0	8	1
1	1	0	3	3
2	0	1	8	2
3	1	0	3	3
4	0	1	8	1
5	0	1	8	1
6	0	1	5	1
7	4	0	8	4
8	2	0	8	3
9	1	0	8	3

```
[55]: X_train = train_df.drop("Survived", axis=1)
Y_train = train_df["Survived"]
X_test = test_df.drop("PassengerId", axis=1).copy()
```

```
[56]: sgd = linear_model.SGDClassifier(max_iter=5, tol=None)
sgd.fit(X_train, Y_train)

Y_pred = sgd.predict(X_test)

sgd.score(X_train, Y_train)
acc_sgd = round(sgd.score(X_train, Y_train) * 100, 2)

# Print score
print(round(acc_sgd,2), "%")
```

58.14 %

```
[57]: decision_tree = DecisionTreeClassifier()
decision_tree.fit(X_train, Y_train)
```

```

Y_pred = decision_tree.predict(X_test)

acc_decision_tree = round(decision_tree.score(X_train, Y_train) * 100, 2)

# Print score
print(round(acc_decision_tree,2,), "%")

```

92.59 %

```

[58]: random_forest = RandomForestClassifier(n_estimators=100)
random_forest.fit(X_train, Y_train)

Y_prediction = random_forest.predict(X_test)

random_forest.score(X_train, Y_train)
acc_random_forest = round(random_forest.score(X_train, Y_train) * 100, 2)

# Print score
print(round(acc_random_forest,2,), "%")

```

92.59 %

```

[59]: logreg = LogisticRegression()
logreg.fit(X_train, Y_train)

Y_pred = logreg.predict(X_test)

acc_log = round(logreg.score(X_train, Y_train) * 100, 2)

# Print score
print(round(acc_log,2,), "%")

```

81.48 %

```

[60]: knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, Y_train)

Y_pred = knn.predict(X_test)

acc_knn = round(knn.score(X_train, Y_train) * 100, 2)

# Print score
print(round(acc_knn,2,), "%")

```

85.52 %

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other

reduction functions (e.g. ``skew``, ``kurtosis``), the default behavior of ``mode`` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of ``keepdims`` will become `False`, the ``axis`` over which the statistic is taken will be eliminated, and the value `None` will no longer be accepted. Set ``keepdims`` to `True` or `False` to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\ProgramData\Anaconda3\lib\site-

packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. ``skew``, ``kurtosis``), the default behavior of ``mode`` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of ``keepdims`` will become `False`, the ``axis`` over which the statistic is taken will be eliminated, and the value `None` will no longer be accepted. Set ``keepdims`` to `True` or `False` to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
[61]: gaussian = GaussianNB()
gaussian.fit(X_train, Y_train)

Y_pred = gaussian.predict(X_test)

acc_gaussian = round(gaussian.score(X_train, Y_train) * 100, 2)

# Print score
print(round(acc_gaussian,2), "%")
```

78.68 %

```
[62]: perceptron = Perceptron(max_iter=1000)
perceptron.fit(X_train, Y_train)

Y_pred = perceptron.predict(X_test)

acc_perceptron = round(perceptron.score(X_train, Y_train) * 100, 2)

# Print score
print(round(acc_perceptron,2), "%")
```

75.76 %

```
[63]: results = pd.DataFrame({
    'Model': ['KNN', 'Logistic Regression',
              'Random Forest', 'Naive Bayes', 'Perceptron',
              'Stochastic Gradient Decent',
              'Decision Tree'],
    'Score': [acc_knn, acc_log,
              acc_random_forest, acc_gaussian, acc_perceptron,
              acc_sgd, acc_decision_tree]})
```

```
result_df = results.sort_values(by='Score', ascending=False)
result_df = result_df.set_index('Score')
result_df.head(9)
```

```
[63]:
```

Score	Model
92.59	Random Forest
92.59	Decision Tree
85.52	KNN
81.48	Logistic Regression
78.68	Naive Bayes
75.76	Perceptron
58.14	Stochastic Gradient Decent

```
[64]: from sklearn.model_selection import cross_val_score

rf = RandomForestClassifier(n_estimators=100)
scores = cross_val_score(rf, X_train, Y_train, cv=10, scoring = "accuracy")
```

```
[65]: print("Scores:", scores)
print("Mean:", scores.mean())
print("Standard Deviation:", scores.std())
```

```
Scores: [0.81111111 0.82022472 0.73033708 0.83146067 0.87640449 0.83146067
 0.79775281 0.7752809 0.85393258 0.82022472]
Mean: 0.8148189762796505
Standard Deviation: 0.03865286486020869
```

```
[66]: importances = pd.DataFrame({'feature':X_train.columns,'importance':np.
    ↳round(random_forest.feature_importances_,3)})
importances = importances.sort_values('importance',ascending=False).
    ↳set_index('feature')
```


```
[67]: importances.head(12)
```

```
[67]:
```

feature	importance
Title	0.198
Sex	0.185
Age	0.147
Fare	0.109
Deck	0.094
Pclass	0.084
relatives	0.067
Embarked	0.047
SibSp	0.033
Parch	0.026

```
not_alone      0.011
train_test     0.000
```

```
[68]: importances.plot.bar();
```



```
[69]: # Dropping not_alone
train_df = train_df.drop("not_alone", axis=1)
test_df  = test_df.drop("not_alone", axis=1)

# Dropping Parch
train_df = train_df.drop("Parch", axis=1)
test_df  = test_df.drop("Parch", axis=1)
```

```
[70]: ## Reassigning features
X_train = train_df.drop("Survived", axis=1)
Y_train = train_df["Survived"]
X_test  = test_df.drop("PassengerId", axis=1).copy()
```



```
[71]: random_forest = RandomForestClassifier(n_estimators=100, oob_score = True)
random_forest.fit(X_train, Y_train)

Y_prediction = random_forest.predict(X_test)

random_forest.score(X_train, Y_train)
acc_random_forest = round(random_forest.score(X_train, Y_train) * 100, 2)

# Print scores
print(round(acc_random_forest,2), "%")
```

92.59 %

```
[72]: importances = pd.DataFrame({'feature':X_train.columns,'importance':np.
    ↪round(random_forest.feature_importances_,3)})
importances = importances.sort_values('importance',ascending=False).
    ↪set_index('feature')
```

```
[73]: importances.head(12)
```

```
[73]:
```

	importance
feature	
Title	0.234
Sex	0.154
Age	0.145
Fare	0.107
Deck	0.094
Pclass	0.090
relatives	0.084
SibSp	0.047
Embarked	0.046
train_test	0.000

```
[74]: print("oob score:", round(random_forest.oob_score_, 4)*100, "%")
```

oob score: 81.47999999999999 %

```
[75]: # Simple performance reporting function
def clf_performance(classifier, model_name):
    print(model_name)
    print('Best Score: ' + str(classifier.best_score_))
    print('Best Parameters: ' + str(classifier.best_params_))
```

```
[76]: from sklearn.model_selection import GridSearchCV

rf = RandomForestClassifier(random_state = 1)
param_grid = {'n_estimators': [400,450,500,550],
```

```

        'criterion':['gini','entropy'],
        'bootstrap': [True],
        'max_depth': [15, 20, 25],
        'max_features': ['auto','sqrt', 10],
        'min_samples_leaf': [2,3],
        'min_samples_split': [2,3]}

clf_rf = GridSearchCV(rf, param_grid = param_grid, cv = 5, verbose = True,
    ↪n_jobs = -1)
best_clf_rf = clf_rf.fit(X_train,Y_train)

# Print score
clf_performance(best_clf_rf,'Random Forest')

```

Fitting 5 folds for each of 288 candidates, totalling 1440 fits

Random Forest

Best Score: 0.8338836231247255

Best Parameters: {'bootstrap': True, 'criterion': 'entropy', 'max_depth': 15, 'max_features': 10, 'min_samples_leaf': 3, 'min_samples_split': 2, 'n_estimators': 400}

```

[77]: random_forest = RandomForestClassifier(criterion = "gini",
                                             max_depth = 20,
                                             max_features='auto',
                                             min_samples_leaf = 3,
                                             min_samples_split = 2,
                                             n_estimators=450,
                                             oob_score=True,
                                             random_state=1,
                                             n_jobs=-1)

random_forest.fit(X_train, Y_train)
Y_prediction = random_forest.predict(X_test)

random_forest.score(X_train, Y_train)

print("oob score:", round(random_forest.oob_score_, 4)*100, "%")

```

oob score: 82.72 %

```

[78]: from sklearn.model_selection import cross_val_predict
      from sklearn.metrics import confusion_matrix

predictions = cross_val_predict(random_forest, X_train, Y_train, cv=3)
confusion_matrix(Y_train, predictions)

```

```
[78]: array([[493,  56],
           [ 97, 245]], dtype=int64)
```

```
[79]: from sklearn.metrics import precision_score, recall_score

print("Precision:", precision_score(Y_train, predictions))
print("Recall:", recall_score(Y_train, predictions))
```

```
Precision: 0.813953488372093
Recall: 0.716374269005848
```

```
[80]: from sklearn.metrics import f1_score
f1_score(Y_train, predictions)
```

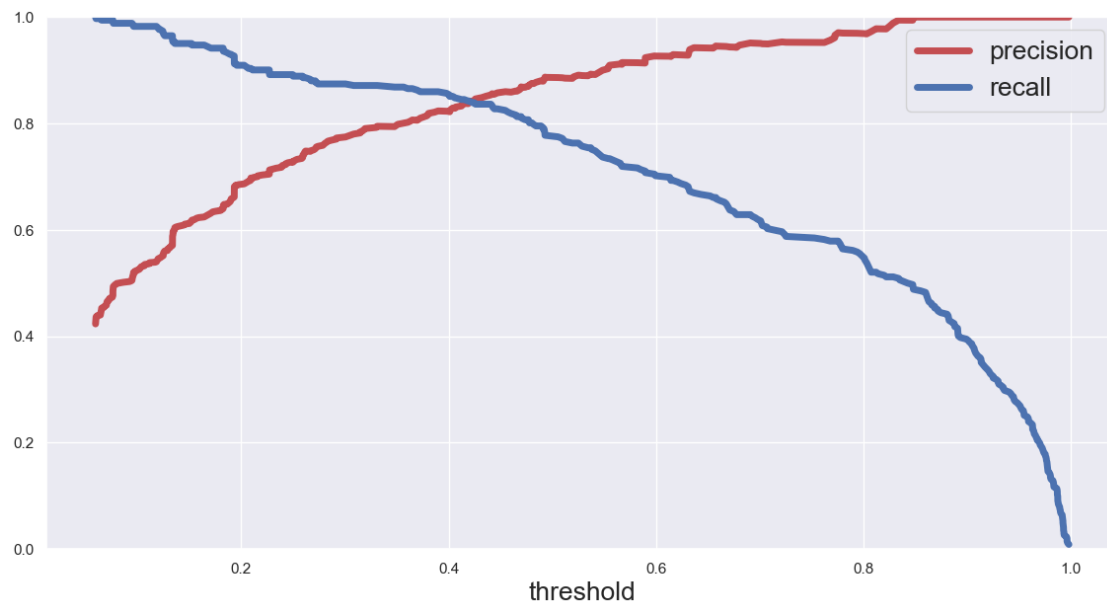
```
[80]: 0.7620528771384137
```

```
[82]: from sklearn.metrics import precision_recall_curve

# Getting the probabilities of our predictions
y_scores = random_forest.predict_proba(X_train)
y_scores = y_scores[:,1]


precision, recall, threshold = precision_recall_curve(Y_train, y_scores)
def plot_precision_and_recall(precision, recall, threshold):
    plt.plot(threshold, precision[:-1], "r", label="precision", linewidth=5)
    plt.plot(threshold, recall[:-1], "b", label="recall", linewidth=5)
    plt.xlabel("threshold", fontsize=19)
    plt.legend(loc="upper right", fontsize=19)
    plt.ylim([0, 1])

plt.figure(figsize=(14, 7))
plot_precision_and_recall(precision, recall, threshold)
plt.show()
```



```
[83]: def plot_precision_vs_recall(precision, recall):
    plt.plot(recall, precision, "b--", linewidth=3)
    plt.xlabel("precision", fontsize=19)
    plt.ylabel("recall", fontsize=19)
    plt.axis([0, 1.2, 0, 1.4])

    plt.figure(figsize=(14, 7))
    plot_precision_vs_recall(precision, recall)
    plt.show()
```

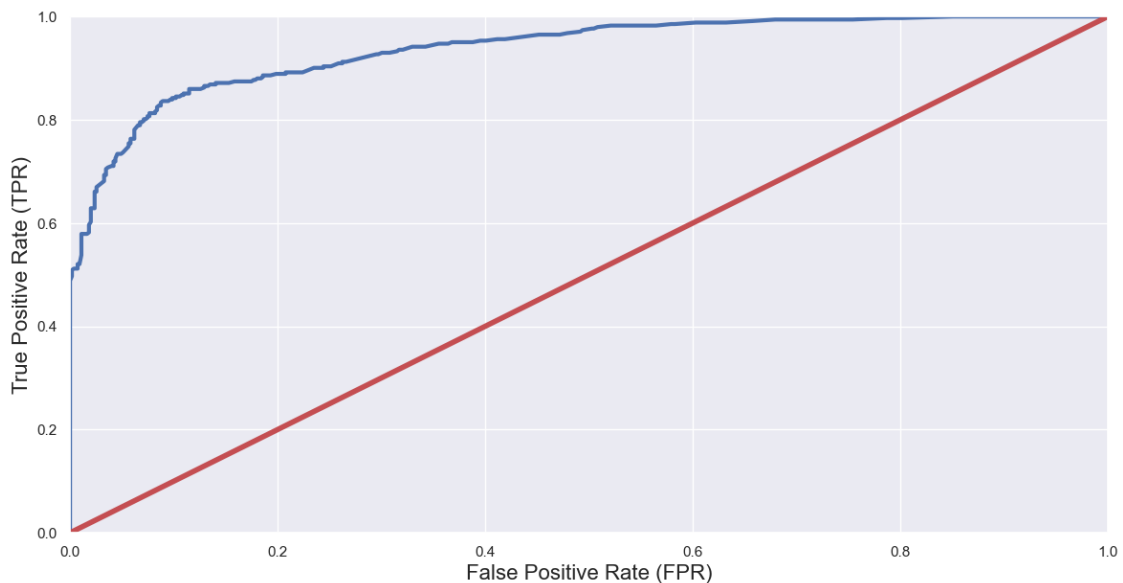


```
[84]: from sklearn.metrics import roc_curve

# Compute true positive rate and false positive rate
false_positive_rate, true_positive_rate, thresholds = roc_curve(Y_train,
    ↪ y_scores)

# Plotting them against each other
def plot_roc_curve(false_positive_rate, true_positive_rate, label=None):
    plt.plot(false_positive_rate, true_positive_rate, linewidth=3, label=label)
    plt.plot([0, 1], [0, 1], 'r', linewidth=4)
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate (FPR)', fontsize=16)
    plt.ylabel('True Positive Rate (TPR)', fontsize=16)

plt.figure(figsize=(14, 7))
plot_roc_curve(false_positive_rate, true_positive_rate)
plt.show()
```



```
[85]: from sklearn.metrics import roc_auc_score
r_a_score = roc_auc_score(Y_train, y_scores)
print("ROC-AUC-Score:", r_a_score)
```

ROC-AUC-Score: 0.9366604885011558

1 ROC AUC SCORE = 93%

[]:

[]: